

# Tutorial: Introduction to Using the Container Storage Interface (CSI) Primitives

*Michael Mattsson, Hewlett Packard Enterprise  
@datamattsson*





# Tutorial: Introduction to Using the Container Storage Interface (CSI) Primitives

*Michael Mattsson, Hewlett Packard Enterprise  
@datamattsson*

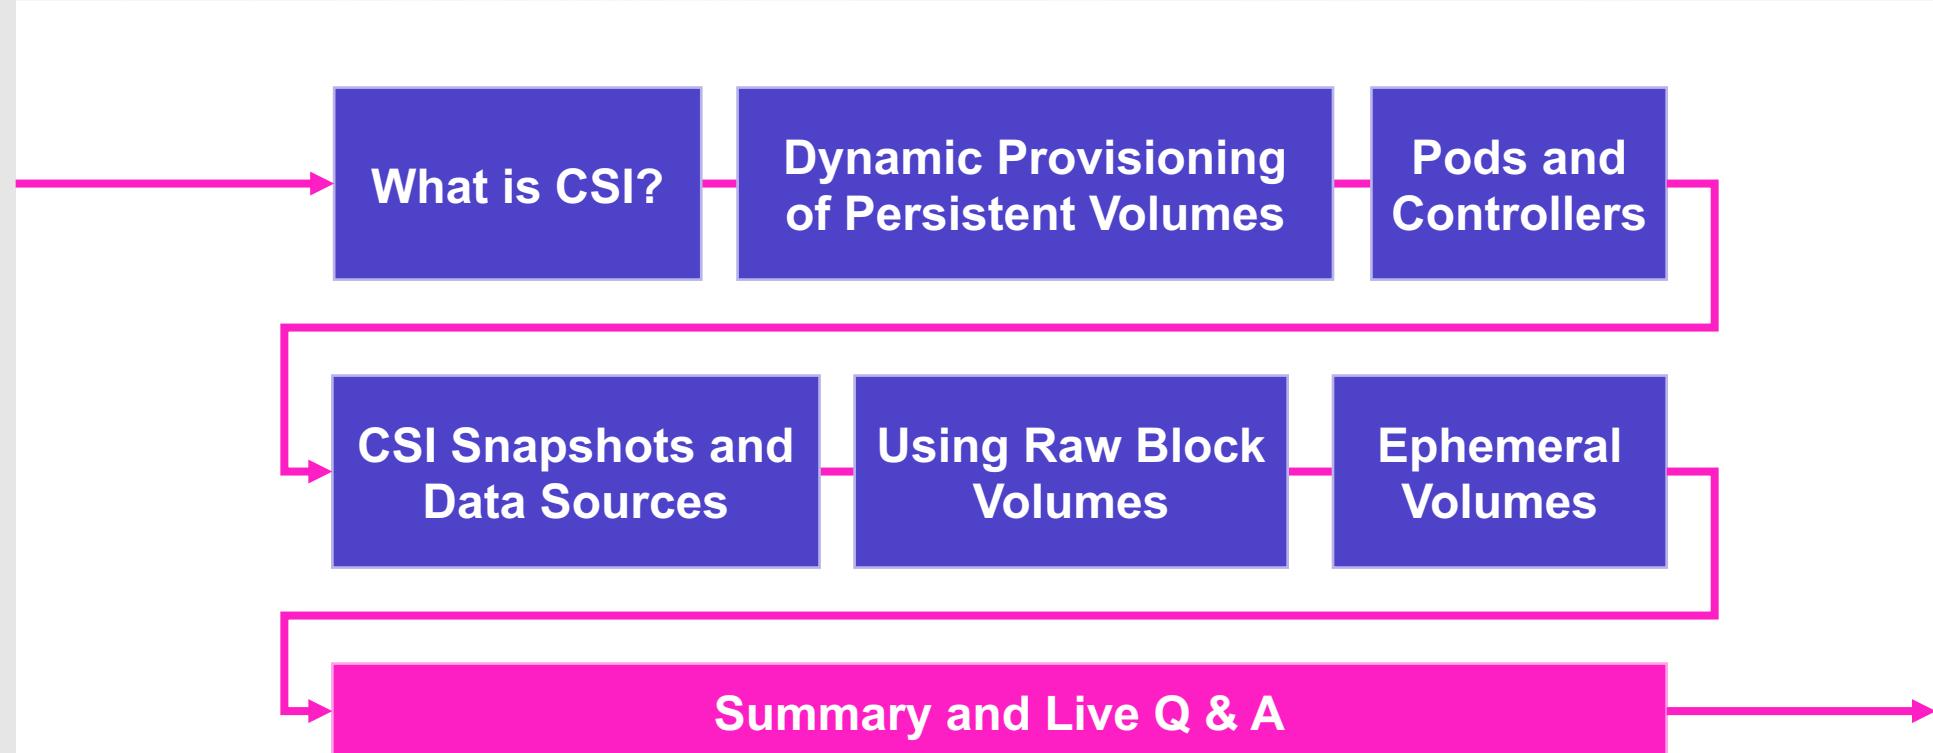
# Tutorial Overview



*Virtual*



CSI stands for  
Container Storage  
Interface



## Shell

```
$ git clone https://github.com/datamattsson/kcna2020
$ cd kcna2020
```

# What is CSI?



# Container Storage Interface



*Virtual*



SIG stands for Special Interest Group



CONTAINER  
STORAGE  
INTERFACE

## CSI Specification

- <https://github.com/container-storage-interface/spec>

## Kubernetes SIG Storage

- <https://github.com/kubernetes/community/tree/master/sig-storage>

## CSI Documentation

- <https://kubernetes-csi.github.io/>

## Container Orchestrators (CO)

- Kubernetes
- Nomad
- Cloud Foundry
- Mesos

# Container Storage Interface

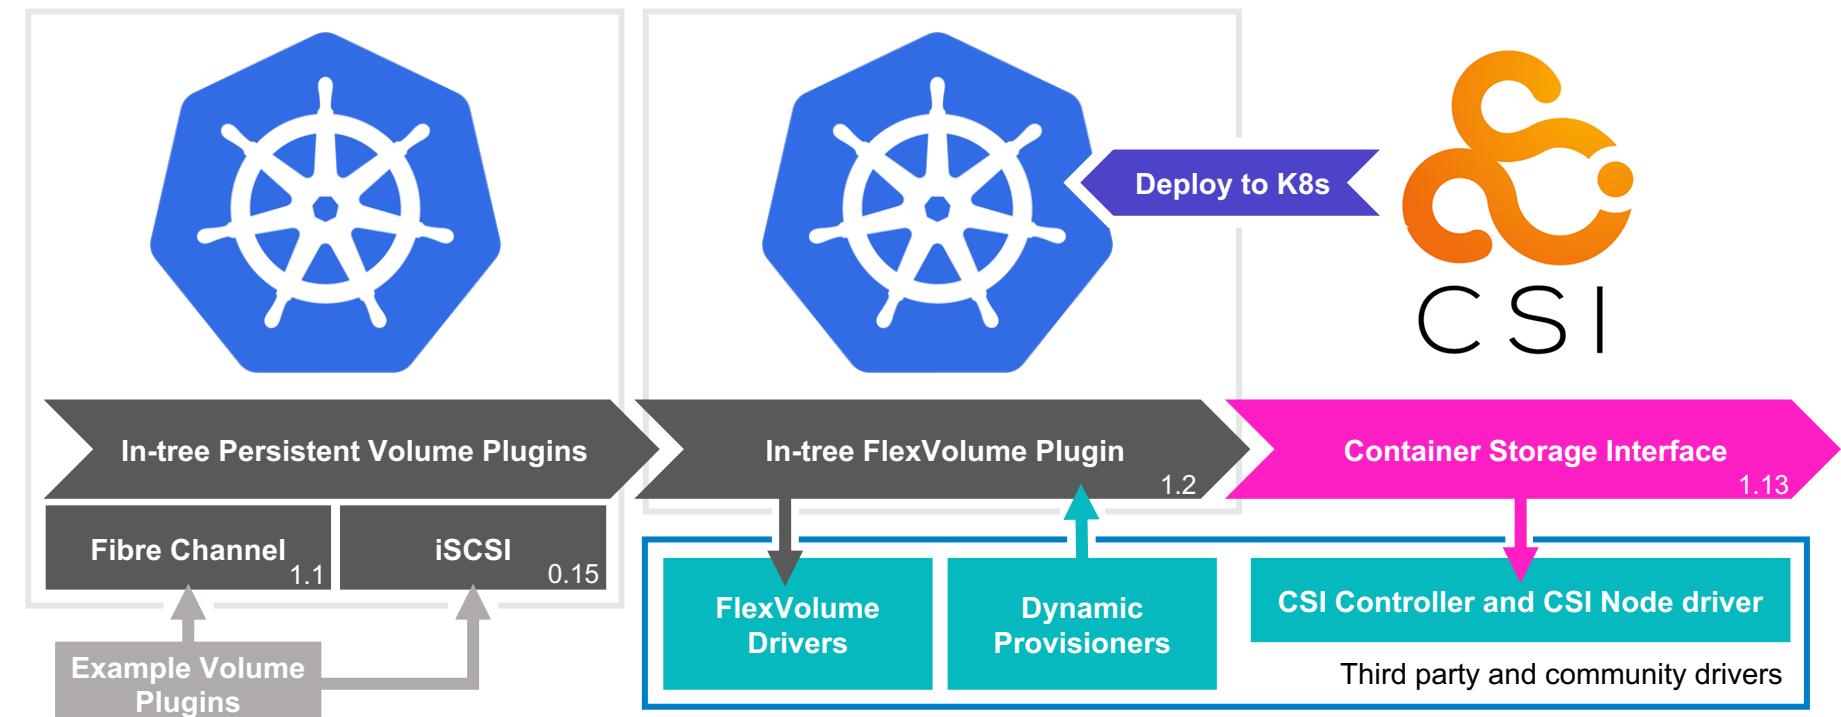


Virtual



In-tree storage plugins  
are being deprecated

## Kubernetes Storage Plugins History



# Container Storage Interface

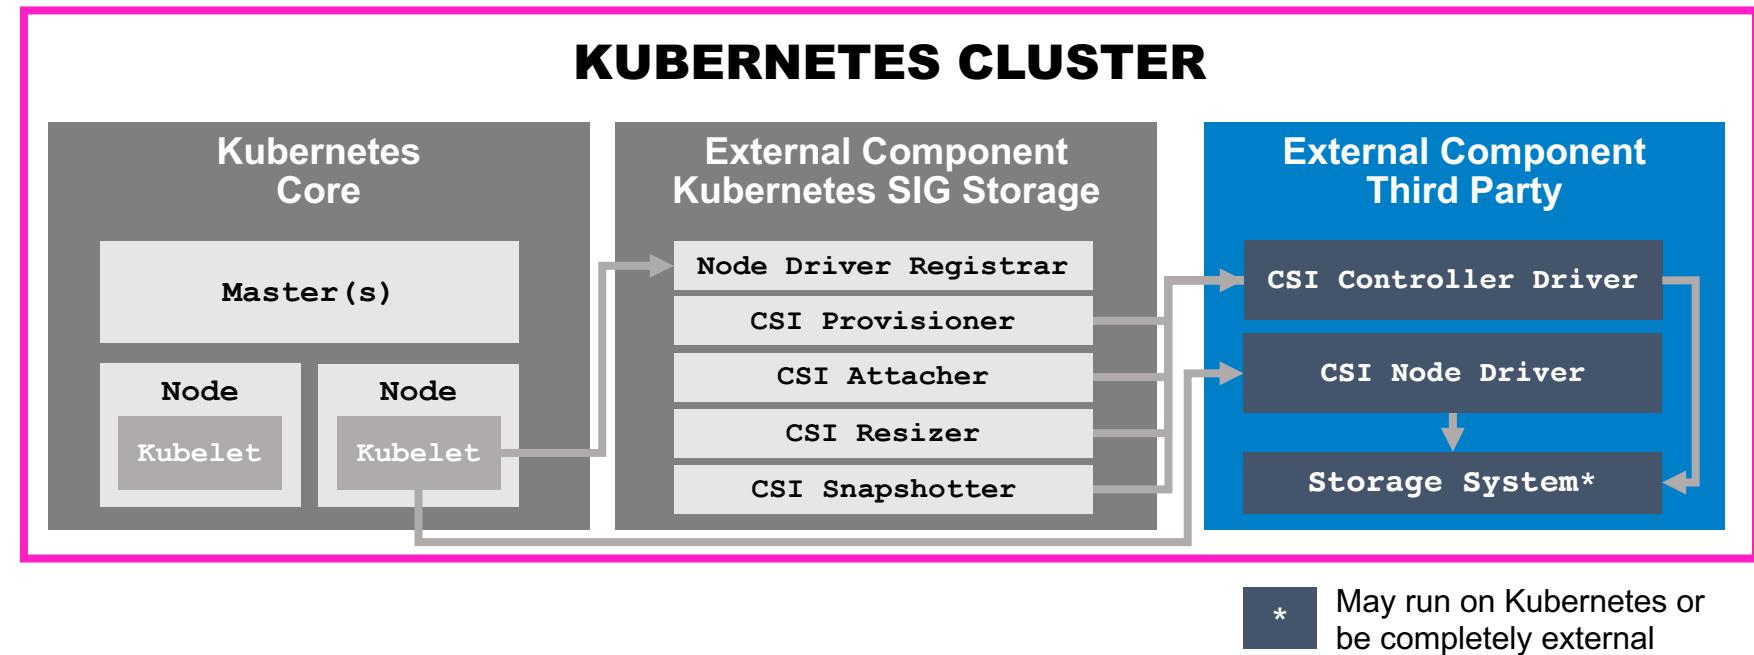


*Virtual*



CSI drivers may provide either file or block storage

## CSI Architecture (simplified)



# Container Storage Interface



*Virtual*



csi.hpe.com will be  
used for the demos

## CSI Drivers

```
diskplugin.csi.alibabacloud.com  csi-vxflexos.dell EMC.com  infinibox-csi-driver  quobyte-csi
nasplugin.csi.alibabacloud.com  csi-xtremio.dell EMC.com  csi-instorage  robin
ossplugin.csi.alibabacloud.com  org.democratic-csi.[X]  pmem-csi.intel.com  csi-sandstone-plugin
arstor.csi.huayun.io  dcx.csi.diamanti.com  csi.juicefs.com  eds.csi.sangfor.com
ebs.csi.aws.com  dobs.csi.digitalocean.com  org.kadalu.gluster  seaweedfs-csi-driver
efs.csi.aws.com  csi.drivescale.com  linodebs.csi.linode.com  secrets-store.csi.k8s.io
fsx.csi.aws.com  [x].ember-csi.io  io.drbd.linstor-csi  csi-smtpx-plugin
disk.csi.azure.com  nvme-csi.excelero.com  driver.longhorn.io  csi.spdk.io
file.csi.azure.com  pd.csi.storage.gke.io  csi-macrosan  storageos
csi.block.bigteracom  com.google.csi.filestore  manila.csi.openstack.org  csi.cio.storidge.com
csi.fs.bigteracom  gcs.csi.ofek.dev  com.mapr.csi-kdf  csi-driver.storpool.com
cephfs.csi.ceph.com  org.gluster.glusterfs  com.tuxera.csi.moosefs  com.tencent.cloud.csi.cbs
rbd.csi.ceph.com  org.gluster.glustervirtblock  csi.trident.netapp.io  com.tencent.cloud.csi.cfs
csi.chubaofs.com  com.hammerspace.csi  nexentastor-csi-driver.nexenta.com  com.tencent.cloud.csi.cosfs
cinder.csi.openstack.org  io.hedvig.csi  nexentastor-block-csi-driver.nexenta.com  topolvm.cybozu.com
csi.cloudscale.ch  csi.hetzner.cloud  com.nutanix.csi  csi.vastdata.com
csi-infiblock-plugin  hspc.csi.hitachi.com  cstor.csi.openebs.io  csi.block.xsky.com
csi-infiflfs-plugin  csi.hpe.com  csi-opensdsplugin  csi.fs.xsky.com
dsp.csi.daterainc.io  csi.huawei.com  com.open-e.joviandss.csi  secrets.csi.kubevault.com
csi-isilon.dell EMC.com  eu.zetanova.csi.hyperv  pxd.openstorage.org  csi.vsphere.vmware.com
csi-powermax.dell EMC.com  block.csi.ibm.com  pure-csi  csi.weka.io
csi-powerstore.dell EMC.com  spectrumscale.csi.ibm.com  disk.csi.qingcloud.com  yandex.csi.flant.com
csi-unity.dell EMC.com  vpc.block.csi.ibm.io  csi-neonsan  csi.zadara.com
```

<https://kubernetes-csi.github.io/docs/drivers.html>

# Container Storage Interface



CloudNativeCon

North America 2020

Virtual



Drivers supports various aspects of the CSI spec

## CSI Driver Example Capabilities

	CephFS	Ceph RDB
<b>Provisioner</b>	cephfs.csi.ceph.com	rbd.csi.ceph.com
<b>Specification compliance</b>	v0.3, v1.0.0, v1.1.0, v1.2.0	v0.3, v1.0.0, v1.1.0, v1.2.0
<b>Modes*</b>	Persistent	Persistent
<b>Access mode</b>	Read/Write Multiple Pods	Read/Write Single Pod
<b>Features</b>	Expansion, Snapshot, Clone	Raw Block, Snapshot, Expansion, Topology, Cloning

\* = May be Persistent or Ephemeral

# Container Storage Interface



CloudNativeCon

North America 2020

Virtual



New features are introduced through Kubernetes Enhancement Proposals (KEP)

## CSI Feature Maturity

Feature	K8s maturity	Since K8s version
<b>Dynamic Provisioning</b> Parameter Overloading	GA	1.13
<b>Topology</b>	GA	1.17
<b>Volume Limits</b>	GA	1.17
<b>Raw Block Volume</b>	GA	1.18
<b>PVC Data Source</b> Cloning	GA	1.18
<b>Volume Expansion</b>	Beta	1.16
<b>Ephemeral Local Volumes</b>	Beta	1.16
<b>Volume Snapshots</b> Volume Snapshot Clone and Restore	Beta	1.17
<b>Generic Ephemeral Volumes</b>	Alpha	1.19

# Container Storage Interface

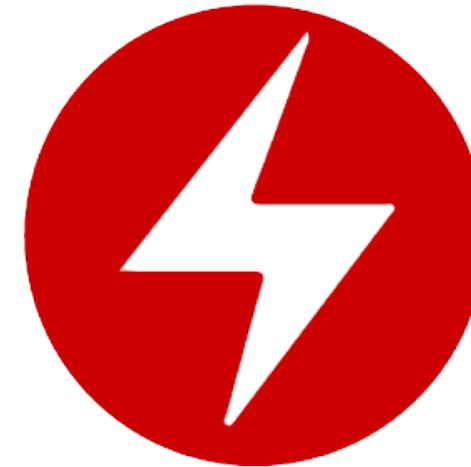


*Virtual*



CSI drivers are usually already installed and setup when using managed Kubernetes

## Installing and inspecting CSI drivers



### Shell

```
$ kubectl get csidrivers  
$ kubectl get csinodes
```

# HANDS-ON #1

Install a CSI driver.

# Dynamic Provisioning of Persistent Volumes

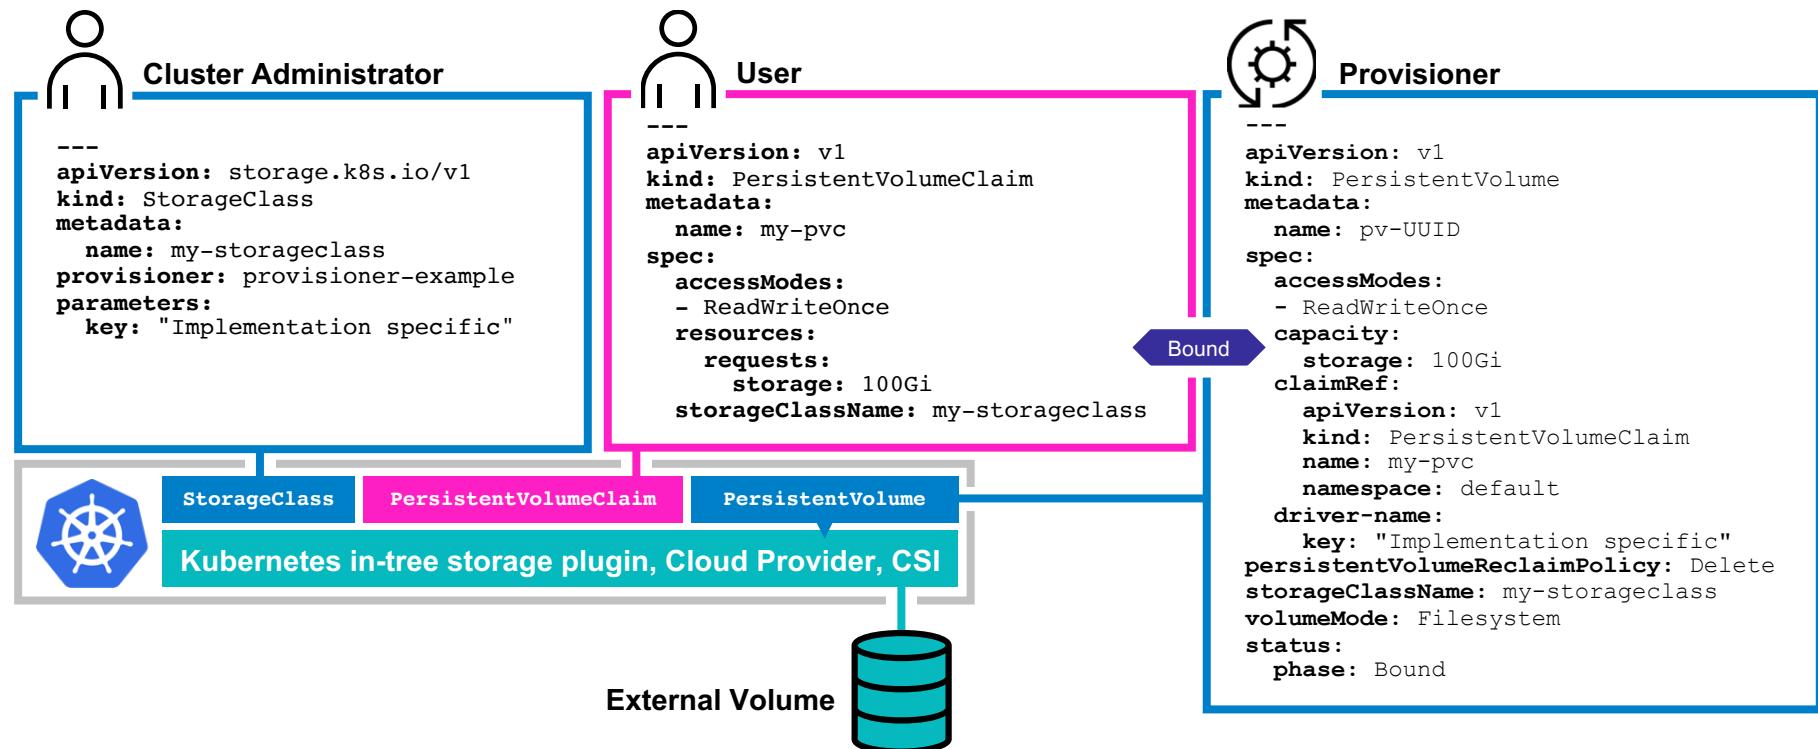


# Dynamic Provisioning



Dynamic provisioning is not exclusive to CSI drivers

## Dynamic Provisioning in Kubernetes



# Dynamic Provisioning



Virtual



StorageClasses are non-namespaced API objects

## StorageClass Overview

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  annotations:
    storageclass.kubernetes.io/is-default-class: "false"
  name: my-storageclass
provisioner: csi.vendor.io
parameters:
  csi.storage.k8s.io/fstype: xfs
  csi.storage.k8s.io/controller-publish-secret-name: csi
  csi.storage.k8s.io/controller-publish-secret-namespace: vendor
...
  csi.storage.k8s.io/controller-expand-secret-name: csi
  csi.storage.k8s.io/controller-expand-secret-namespace: vendor
reclaimPolicy: Delete
volumeBindingMode: Immediate
allowVolumeExpansion: "false"
```

- Optional, with their respective defaults
- ... Each CSI Sidecar needs to reference a secret

# Dynamic Provisioning



Virtual



A PVC is confined to a Namespace

## Persistent Volume Claim Overview

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-pvc
spec:
  accessModes:
  - ReadWriteOnce
  resources:
    requests:
      storage: 2Ti
  volumeMode: Filesystem
  storageClassName: my-storageclass
```

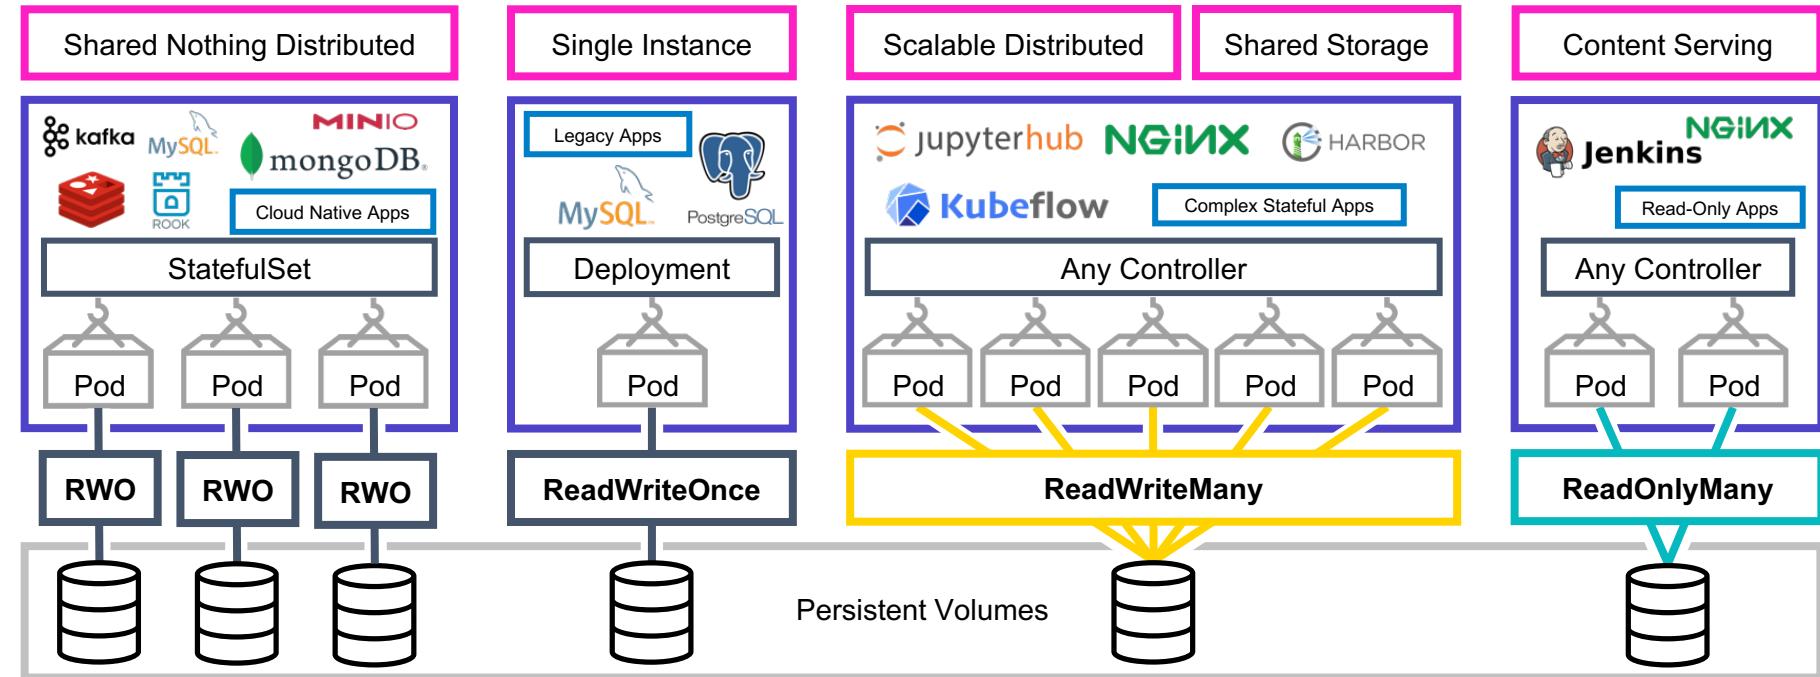
- Optional, with their respective defaults
- Uses the default **storageClass** if omitted

# Dynamic Provisioning



A PVC is confined to a Namespace

## PVC Access Modes



# Dynamic Provisioning



KubeCon



CloudNativeCor

## — North America 2020

# Virtual



The **PV** describes the storage resource in a way the CSI driver understand

# Persistent Volume Overview

... Each CSI operation  
needs to reference a secret

Note: `claimRef` omitted

# HANDS-ON #2

Create a StorageClass and PVC. Attach a workload and expand volume.

# Dynamic Provisioning



All workload controllers except the StatefulSet works very similar

## Workload Controllers

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
    - name: myfrontend
      image: nginx
      volumeMounts:
        - mountPath: "/usr/share/nginx/html"
          name: my-mount
  volumes:
    - name: my-mount
      persistentVolumeClaim:
        claimName: my-pvc
```

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: my-web
spec:
  selector:
    matchLabels:
      app: nginx
  serviceName: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
          volumeMounts:
            - name: my-mount
              mountPath: "/usr/share/nginx/html"
  volumeClaimTemplates:
  - metadata:
      name: my-mount
    spec:
      accessModes: [ "ReadWriteOnce" ]
      storageClassName: my-storageclass
      resources:
        requests:
          storage: 2Ti
```

# HANDS-ON #3

Deploy an application utilizing a StatefulSet.

# CSI Snapshots and PVC Data Sources



# CSI Snapshot and dataSource

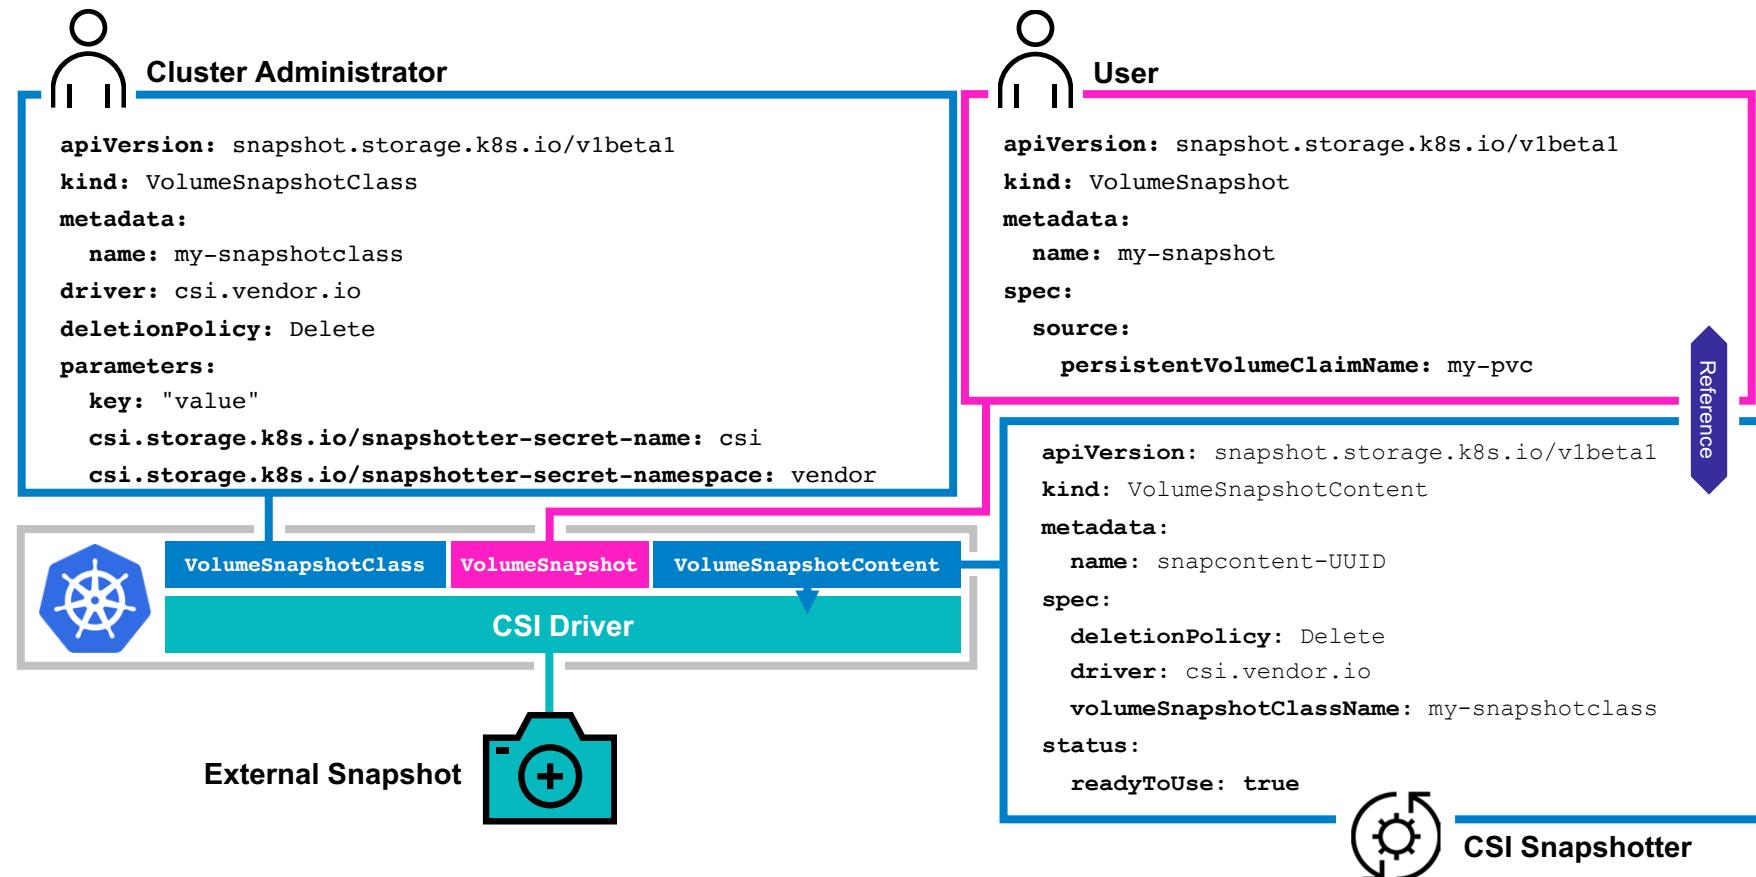


Virtual



Snapshots uses storage system snapshots abstracted to API objects

## CSI Snapshots



# HANDS-ON #4

Deploy CSI snapshotter, create a VolumeSnapshotClass and VolumeSnapshots.



New PVCs may be created from VolumeSnapshots

## PVC from VolumeSnapshot

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-new-pvc
spec:
  dataSource:
    name: my-snapshot
    kind: VolumeSnapshot
    apiGroup: snapshot.storage.k8s.io
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 2Ti
```

# HANDS-ON #5

Create a new PVC from a VolumeSnapshot and attach an application.



A PVC may be created from an existing PVC

## PVC Cloning

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-clone-pvc
spec:
  dataSource:
    name: my-pvc
    kind: PersistentVolumeClaim
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 2Ti
```

# HANDS-ON #6

Create a new PVC from an existing PVC and attach an application.



Restoring an application  
from a `VolumeSnapshot`  
is a cloning operation

## Restore an application

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-pvc
spec:
  dataSource:
    name: my-snapshot
    kind: VolumeSnapshot
    apiGroup: snapshot.storage.k8s.io
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 2Ti
```

# HANDS-ON #7

Restore an application from a VolumeSnapshot.

# Using Raw Block Volumes



# Using Raw Block Volumes



Virtual



The default `volumeMode` is filesystem

## Raw Block Volume PVC

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-block-pvc
spec:
  accessModes:
  - ReadWriteOnce
  resources:
    requests:
      storage: 2Ti
  volumeMode: Block
  storageClassName: my-storageclass
```

- Uses the default `storageClass` if omitted

# Using Raw Block Volumes



*Virtual*



Devices are referenced  
and enumerated slightly  
different from filesystems

## Raw Block Volume Pod specification

```
apiVersion: v1
kind: Pod
metadata:
  name: ioping
spec:
  containers:
  - name: ioping
    image: datamattsson/ioping:edge
    args: [ "/dev/xvda" ]
    volumeDevices:
    - name: volume
      devicePath: /dev/xvda
  volumes:
  - name: volume
    persistentVolumeClaim:
      claimName: my-block-pvc
```

# HANDS-ON #8

Create a raw block device and attach a workload.

# Using Raw Block Volumes



Virtual



Rook is a CNCF incubating project

## Use Case for Raw Block Volumes



### Rook

- Open-Source, Cloud-Native Storage for Kubernetes
- Provides File, Block and Object to Kubernetes
- Uses CEPH
- Manage distributed storage on Kubernetes effortless

### Deployment pattern

- Installs as an Operator
- Create **CephCluster** CRD
- May leverage **volumeMode: Block** PVC

# HANDS-ON #9

Deploy Rook.

# Ephemeral Volumes



# Ephemeral Volumes



Ephemeral Volumes exist only for the duration of the **Pod** and inherently **ReadWriteOnce** per **Pod**



*Virtual*

## Ephemeral Local Volumes

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
    - name: myfrontend
      image: nginx
      volumeMounts:
        - mountPath: "/usr/share/nginx/html"
          name: my-mount
  volumes:
    - name: my-mount
      csi:
        driver: csi.vendor.io
        volumeAttributes:
          csi.storage.k8s.io/ephemeral: "true"
          inline-volume-secret-name: csi
          inline-volume-secret-namespace: vendor
          size: 64Gi
```

# Ephemeral Volumes



Ephemeral Volumes exist only for the duration of the Pod



*Virtual*

## Ephemeral Local Volumes (alternative `secretRef`)

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
    - name: myfrontend
      image: nginx
      volumeMounts:
        - mountPath: "/usr/share/nginx/html"
          name: my-mount
  volumes:
    - name: my-mount
      csi:
        driver: csi.vendor.io
        volumeAttributes:
          csi.storage.k8s.io/ephemeral: "true"
        nodePublishSecretRef:
          name: csi
        size: 64Gi
```

# **HANDS-ON #10**

Using Ephemeral Local Volumes.

# Ephemeral Volumes



Generic Ephemeral Volumes is an Alpha feature introduced in 1.19



Virtual

## Generic Ephemeral Volumes

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
    - name: myfrontend
      image: nginx
      volumeMounts:
        - mountPath: "/usr/share/nginx/html"
          name: my-mount
  volumes:
    - name: my-mount
      ephemeral:
        volumeClaimTemplate:
          spec:
            accessModes: [ "ReadWriteOnce" ]
            resources:
              requests:
                storage: 64Gi
            storageClassName: my-storageclass
```

- Uses the default **storageClass** if omitted

# **HANDS-ON #11**

Using Generic Ephemeral Volumes.

# Summary



*Virtual*



We are done!

## Thank you for participating!



CONTAINER  
STORAGE  
INTERFACE

### What we covered

- Introduction to CSI drivers
- Dynamic provisioning in Kubernetes
- CSI snapshots, restore and PVC cloning
- Accessing raw block storage
- Ephemeral volumes

### Source files (YAML, .pptx and asciinema cast files)

- <https://github.com/datamattsson/kcna2020>

### CSI Specification

- <https://github.com/container-storage-interface/spec>

### Kubernetes SIG Storage

- <https://github.com/kubernetes/community/tree/master/sig-storage>

### CSI Documentation

- <https://kubernetes-csi.github.io/>

